

```
In [ ]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import os
import pickle
import matplotlib.patches as patches
import re
import random
from sklearn.model_selection import train_test_split
import cv2
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
from keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.utils import plot_model
from PIL import Image
import tensorflow as tf
from keras.layers import GlobalAveragePooling2D,Dense,Conv2D,BatchNormalization,Dropout
import keras
from keras import backend as K
from keras.models import Model,load_model
from tensorflow.python.framework.ops import disable_eager_execution
from keras.regularizers import l2
import datetime
%load_ext tensorboard
```

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [ ]: with open('/content/drive//My Drive/Steel_Detection /data.pkl','rb') as f:
        train=pickle.load(f)
```

```
In [ ]: train.head()
```

	image_id	rle_1	rle_2	rle_3	rle_4	defect	stratify	defect_1	defect_2	defect_3	defect_4	total_defects
0	0002cc93b.jpg	29102 12 29346 24 29602 24 29858 24 30114 24 3...				1	1	1	0	0	0	1
1	00031f466.jpg					0	0	0	0	0	0	0
2	000418bfc.jpg					0	0	0	0	0	0	0
3	000789191.jpg					0	0	0	0	0	0	0
4	0007a71bf.jpg	18661 28 18863 82 19091 110 19347 110 19603 11...				1	3	0	0	1	0	1

```
In [ ]: train.shape
```

```
(12568, 12)
```

```
In [ ]: def f1_score(y_true, y_pred):
    #https://aakashgoel12.medium.com/how-to-add-user-defined-function-get-f1-score-in-keras-metric-s-3013f979ce0d
    #https://stackoverflow.com/questions/43547402/how-to-calculate-f1-macro-in-keras
    true_positives=K.sum(K.round(K.clip(y_true*y_pred,0,1)))    #calculates number of true positive
    possible_positives=K.sum(K.round(K.clip(y_true,0,1)))        #calculates number of actual positives
    predicted_positives=K.sum(K.round(K.clip(y_pred,0,1)))

    #K.epsilon takes care of non-zero divisions
    #was modified by adding the constant epsilon, in order to avoid division by 0. Thus NaN will not be computed.
    precision=true_positives/(predicted_positives +K.epsilon())
    recall=true_positives/(possible_positives+K.epsilon())
    f1_val=2*(precision*recall)/(precision+recall+K.epsilon())
    return f1_val
```

- If image have defect than only image considered in data.

```
In [ ]: #https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html
train=train[train['defect']==1]
x_train,x_test=train_test_split(train,test_size=0.10,stratify=train['stratify'],random_state=0)
x_train,x_val=train_test_split(x_train,test_size=0.20,stratify=x_train['stratify'],random_state=0)
x_train=x_train[['image_id','defect_1','defect_2','defect_3','defect_4']]
x_val=x_val[['image_id','defect_1','defect_2','defect_3','defect_4']]
x_test=x_test[['image_id','defect_1','defect_2','defect_3','defect_4']]
print("x_train {}".format(x_train.shape)," x_val {}".format(x_val.shape)," x_test {}".format(x_test.shape))

x_train (4799, 5)  x_val (1200, 5)  x_test (667, 5)
```

```
In [ ]: #https://www.analyticsvidhya.com/blog/2020/08/image-augmentation-on-the-fly-using-keras-imagedatagenerator/
#https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator
train_datagen=ImageDataGenerator(rescale=1./255,
                                  shear_range=0.2,
                                  zoom_range=0.1,
                                  horizontal_flip=True,
                                  vertical_flip=True,
                                  rotation_range=10,
                                  width_shift_range=0.1,
                                  height_shift_range=0.1)

val_datagen=ImageDataGenerator(rescale=1./255)

train_folder_path='/content/drive//My Drive/Steel_Detection /train_images'
col=["defect_1","defect_2","defect_3","defect_4"]
train_image_generator=train_datagen.flow_from_dataframe(dataframe=x_train,
                                                         directory=train_folder_path,
                                                         x_col="image_id",
                                                         y_col=col,
                                                         batch_size=32,
                                                         class_mode="raw",
                                                         target_size=(256,512))

val_image_generator=val_datagen.flow_from_dataframe(dataframe=x_val,
                                                     directory=train_folder_path,
                                                     x_col="image_id",
                                                     y_col=col,
                                                     batch_size=32,
                                                     class_mode="raw",
                                                     target_size=(256,512))
```

Found 4799 validated image filenames.

Found 1200 validated image filenames.

- Used pre-trained Xception() model as base model by keras without fully-connected layer at the top of the network and later freeze the pre-trained model.

```
In [ ]: #https://keras.io/api/applications/inceptionresnetv2/
base_model=tf.keras.applications.xception.Xception(input_shape=(256,512,3),include_top=False)
base_model.trainable=False

m=base_model.output
# add a global average pooling layer
#https://stackoverflow.com/questions/49295311/what-is-the-difference-between-flatten-and-globalaveragepooling2d-in-keras
m=GlobalAveragePooling2D()(m)

# add fully-connected layers
m=Dense(1024,activation='relu')(m)
m=BatchNormalization()(m) #https://www.analyticsvidhya.com/blog/2021/03/introduction-to-batch-normalization/
m=Dropout(0.4)(m)

#https://machinelearningmastery.com/how-to-reduce-overfitting-in-deep-learning-with-weight-regularization/
m=Dense(512,activation='relu')(m)
m=BatchNormalization()(m)
m=Dropout(0.4)(m)

m=Dense(64,activation='relu')(m)

#prediction layer
output=Dense(4,activation='sigmoid')(m)

model=Model(inputs=base_model.input,outputs=output)
model._name="Multi_label_Classification_Model"
model.summary()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/xception/xception_weights_tf_dim_ordering_tf_kernels_no_top.h5

83689472/83683744 [=====] - 1s 0us/step

83697664/83683744 [=====] - 1s 0us/step

Model: "Multi_label_Classification_Model"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	[(None, 256, 512, 3)] 0		
block1_conv1 (Conv2D)	(None, 127, 255, 32) 864		input_1[0][0]
block1_conv1_bn (BatchNormaliza	(None, 127, 255, 32) 128		block1_conv1[0][0]
block1_conv1_act (Activation)	(None, 127, 255, 32) 0		block1_conv1_bn[0][0]
block1_conv2 (Conv2D)	(None, 125, 253, 64) 18432		block1_conv1_act[0][0]
block1_conv2_bn (BatchNormaliza	(None, 125, 253, 64) 256		block1_conv2[0][0]
block1_conv2_act (Activation)	(None, 125, 253, 64) 0		block1_conv2_bn[0][0]
block2_sepconv1 (SeparableConv2	(None, 125, 253, 128) 8768		block1_conv2_act[0][0]
block2_sepconv1_bn (BatchNormal	(None, 125, 253, 128) 512		block2_sepconv1[0][0]
block2_sepconv2_act (Activation	(None, 125, 253, 128) 0		block2_sepconv1_bn[0][0]
block2_sepconv2 (SeparableConv2	(None, 125, 253, 128) 17536		block2_sepconv2_act[0][0]
block2_sepconv2_bn (BatchNormal	(None, 125, 253, 128) 512		block2_sepconv2[0][0]
conv2d (Conv2D)	(None, 63, 127, 128) 8192		block1_conv2_act[0][0]
block2_pool (MaxPooling2D)	(None, 63, 127, 128) 0		block2_sepconv2_bn[0][0]
batch_normalization (BatchNorma	(None, 63, 127, 128) 512		conv2d[0][0]
add (Add)	(None, 63, 127, 128) 0		block2_pool[0][0] batch_normalization[0][0]
block3_sepconv1_act (Activation	(None, 63, 127, 128) 0		add[0][0]

block3_sepconv1 (SeparableConv2 (None, 63, 127, 256) 33920	block3_sepconv1_act[0][0]
block3_sepconv1_bn (BatchNormal (None, 63, 127, 256) 1024	block3_sepconv1[0][0]
block3_sepconv2_act (Activation (None, 63, 127, 256) 0	block3_sepconv1_bn[0][0]
block3_sepconv2 (SeparableConv2 (None, 63, 127, 256) 67840	block3_sepconv2_act[0][0]
block3_sepconv2_bn (BatchNormal (None, 63, 127, 256) 1024	block3_sepconv2[0][0]
conv2d_1 (Conv2D) (None, 32, 64, 256) 32768	add[0][0]
block3_pool (MaxPooling2D) (None, 32, 64, 256) 0	block3_sepconv2_bn[0][0]
batch_normalization_1 (BatchNor (None, 32, 64, 256) 1024	conv2d_1[0][0]
add_1 (Add) (None, 32, 64, 256) 0	block3_pool[0][0] batch_normalization_1[0][0]
block4_sepconv1_act (Activation (None, 32, 64, 256) 0	add_1[0][0]
block4_sepconv1 (SeparableConv2 (None, 32, 64, 728) 188672	block4_sepconv1_act[0][0]
block4_sepconv1_bn (BatchNormal (None, 32, 64, 728) 2912	block4_sepconv1[0][0]
block4_sepconv2_act (Activation (None, 32, 64, 728) 0	block4_sepconv1_bn[0][0]
block4_sepconv2 (SeparableConv2 (None, 32, 64, 728) 536536	block4_sepconv2_act[0][0]
block4_sepconv2_bn (BatchNormal (None, 32, 64, 728) 2912	block4_sepconv2[0][0]
conv2d_2 (Conv2D) (None, 16, 32, 728) 186368	add_1[0][0]
block4_pool (MaxPooling2D) (None, 16, 32, 728) 0	block4_sepconv2_bn[0][0]
batch_normalization_2 (BatchNor (None, 16, 32, 728) 2912	conv2d_2[0][0]
add_2 (Add) (None, 16, 32, 728) 0	block4_pool[0][0] batch_normalization_2[0][0]
block5_sepconv1_act (Activation (None, 16, 32, 728) 0	add_2[0][0]
block5_sepconv1 (SeparableConv2 (None, 16, 32, 728) 536536	block5_sepconv1_act[0][0]

block5_sepconv1_bn	(BatchNormal (None, 16, 32, 728)	2912	block5_sepconv1[0][0]
block5_sepconv2_act	(Activation (None, 16, 32, 728)	0	block5_sepconv1_bn[0][0]
block5_sepconv2	(SeparableConv2 (None, 16, 32, 728)	536536	block5_sepconv2_act[0][0]
block5_sepconv2_bn	(BatchNormal (None, 16, 32, 728)	2912	block5_sepconv2[0][0]
block5_sepconv3_act	(Activation (None, 16, 32, 728)	0	block5_sepconv2_bn[0][0]
block5_sepconv3	(SeparableConv2 (None, 16, 32, 728)	536536	block5_sepconv3_act[0][0]
block5_sepconv3_bn	(BatchNormal (None, 16, 32, 728)	2912	block5_sepconv3[0][0]
add_3 (Add)	(None, 16, 32, 728)	0	block5_sepconv3_bn[0][0] add_2[0][0]
block6_sepconv1_act	(Activation (None, 16, 32, 728)	0	add_3[0][0]
block6_sepconv1	(SeparableConv2 (None, 16, 32, 728)	536536	block6_sepconv1_act[0][0]
block6_sepconv1_bn	(BatchNormal (None, 16, 32, 728)	2912	block6_sepconv1[0][0]
block6_sepconv2_act	(Activation (None, 16, 32, 728)	0	block6_sepconv1_bn[0][0]
block6_sepconv2	(SeparableConv2 (None, 16, 32, 728)	536536	block6_sepconv2_act[0][0]
block6_sepconv2_bn	(BatchNormal (None, 16, 32, 728)	2912	block6_sepconv2[0][0]
block6_sepconv3_act	(Activation (None, 16, 32, 728)	0	block6_sepconv2_bn[0][0]
block6_sepconv3	(SeparableConv2 (None, 16, 32, 728)	536536	block6_sepconv3_act[0][0]
block6_sepconv3_bn	(BatchNormal (None, 16, 32, 728)	2912	block6_sepconv3[0][0]
add_4 (Add)	(None, 16, 32, 728)	0	block6_sepconv3_bn[0][0] add_3[0][0]
block7_sepconv1_act	(Activation (None, 16, 32, 728)	0	add_4[0][0]
block7_sepconv1	(SeparableConv2 (None, 16, 32, 728)	536536	block7_sepconv1_act[0][0]

block7_sepconv1_bn (BatchNormal (None, 16, 32, 728)	2912	block7_sepconv1[0][0]
block7_sepconv2_act (Activation (None, 16, 32, 728)	0	block7_sepconv1_bn[0][0]
block7_sepconv2 (SeparableConv2 (None, 16, 32, 728)	536536	block7_sepconv2_act[0][0]
block7_sepconv2_bn (BatchNormal (None, 16, 32, 728)	2912	block7_sepconv2[0][0]
block7_sepconv3_act (Activation (None, 16, 32, 728)	0	block7_sepconv2_bn[0][0]
block7_sepconv3 (SeparableConv2 (None, 16, 32, 728)	536536	block7_sepconv3_act[0][0]
block7_sepconv3_bn (BatchNormal (None, 16, 32, 728)	2912	block7_sepconv3[0][0]
add_5 (Add) (None, 16, 32, 728)	0	block7_sepconv3_bn[0][0] add_4[0][0]
block8_sepconv1_act (Activation (None, 16, 32, 728)	0	add_5[0][0]
block8_sepconv1 (SeparableConv2 (None, 16, 32, 728)	536536	block8_sepconv1_act[0][0]
block8_sepconv1_bn (BatchNormal (None, 16, 32, 728)	2912	block8_sepconv1[0][0]
block8_sepconv2_act (Activation (None, 16, 32, 728)	0	block8_sepconv1_bn[0][0]
block8_sepconv2 (SeparableConv2 (None, 16, 32, 728)	536536	block8_sepconv2_act[0][0]
block8_sepconv2_bn (BatchNormal (None, 16, 32, 728)	2912	block8_sepconv2[0][0]
block8_sepconv3_act (Activation (None, 16, 32, 728)	0	block8_sepconv2_bn[0][0]
block8_sepconv3 (SeparableConv2 (None, 16, 32, 728)	536536	block8_sepconv3_act[0][0]
block8_sepconv3_bn (BatchNormal (None, 16, 32, 728)	2912	block8_sepconv3[0][0]
add_6 (Add) (None, 16, 32, 728)	0	block8_sepconv3_bn[0][0] add_5[0][0]
block9_sepconv1_act (Activation (None, 16, 32, 728)	0	add_6[0][0]
block9_sepconv1 (SeparableConv2 (None, 16, 32, 728)	536536	block9_sepconv1_act[0][0]
block9_sepconv1_bn (BatchNormal (None, 16, 32, 728)	2912	block9_sepconv1[0][0]

block9_sepconv2_act (Activation (None, 16, 32, 728) 0	block9_sepconv1_bn[0][0]
block9_sepconv2 (SeparableConv2 (None, 16, 32, 728) 536536	block9_sepconv2_act[0][0]
block9_sepconv2_bn (BatchNormal (None, 16, 32, 728) 2912	block9_sepconv2[0][0]
block9_sepconv3_act (Activation (None, 16, 32, 728) 0	block9_sepconv2_bn[0][0]
block9_sepconv3 (SeparableConv2 (None, 16, 32, 728) 536536	block9_sepconv3_act[0][0]
block9_sepconv3_bn (BatchNormal (None, 16, 32, 728) 2912	block9_sepconv3[0][0]
add_7 (Add) (None, 16, 32, 728) 0	block9_sepconv3_bn[0][0] add_6[0][0]
block10_sepconv1_act (Activatio (None, 16, 32, 728) 0	add_7[0][0]
block10_sepconv1 (SeparableConv (None, 16, 32, 728) 536536	block10_sepconv1_act[0][0]
block10_sepconv1_bn (BatchNorma (None, 16, 32, 728) 2912	block10_sepconv1[0][0]
block10_sepconv2_act (Activatio (None, 16, 32, 728) 0	block10_sepconv1_bn[0][0]
block10_sepconv2 (SeparableConv (None, 16, 32, 728) 536536	block10_sepconv2_act[0][0]
block10_sepconv2_bn (BatchNorma (None, 16, 32, 728) 2912	block10_sepconv2[0][0]
block10_sepconv3_act (Activatio (None, 16, 32, 728) 0	block10_sepconv2_bn[0][0]
block10_sepconv3 (SeparableConv (None, 16, 32, 728) 536536	block10_sepconv3_act[0][0]
block10_sepconv3_bn (BatchNorma (None, 16, 32, 728) 2912	block10_sepconv3[0][0]
add_8 (Add) (None, 16, 32, 728) 0	block10_sepconv3_bn[0][0] add_7[0][0]
block11_sepconv1_act (Activatio (None, 16, 32, 728) 0	add_8[0][0]
block11_sepconv1 (SeparableConv (None, 16, 32, 728) 536536	block11_sepconv1_act[0][0]
block11_sepconv1_bn (BatchNorma (None, 16, 32, 728) 2912	block11_sepconv1[0][0]

block11_sepconv2_act (Activatio	(None, 16, 32, 728)	0	block11_sepconv1_bn[0][0]
block11_sepconv2 (SeparableConv	(None, 16, 32, 728)	536536	block11_sepconv2_act[0][0]
block11_sepconv2_bn (BatchNorma	(None, 16, 32, 728)	2912	block11_sepconv2[0][0]
block11_sepconv3_act (Activatio	(None, 16, 32, 728)	0	block11_sepconv2_bn[0][0]
block11_sepconv3 (SeparableConv	(None, 16, 32, 728)	536536	block11_sepconv3_act[0][0]
block11_sepconv3_bn (BatchNorma	(None, 16, 32, 728)	2912	block11_sepconv3[0][0]
add_9 (Add)	(None, 16, 32, 728)	0	block11_sepconv3_bn[0][0] add_8[0][0]
block12_sepconv1_act (Activatio	(None, 16, 32, 728)	0	add_9[0][0]
block12_sepconv1 (SeparableConv	(None, 16, 32, 728)	536536	block12_sepconv1_act[0][0]
block12_sepconv1_bn (BatchNorma	(None, 16, 32, 728)	2912	block12_sepconv1[0][0]
block12_sepconv2_act (Activatio	(None, 16, 32, 728)	0	block12_sepconv1_bn[0][0]
block12_sepconv2 (SeparableConv	(None, 16, 32, 728)	536536	block12_sepconv2_act[0][0]
block12_sepconv2_bn (BatchNorma	(None, 16, 32, 728)	2912	block12_sepconv2[0][0]
block12_sepconv3_act (Activatio	(None, 16, 32, 728)	0	block12_sepconv2_bn[0][0]
block12_sepconv3 (SeparableConv	(None, 16, 32, 728)	536536	block12_sepconv3_act[0][0]
block12_sepconv3_bn (BatchNorma	(None, 16, 32, 728)	2912	block12_sepconv3[0][0]
add_10 (Add)	(None, 16, 32, 728)	0	block12_sepconv3_bn[0][0] add_9[0][0]
block13_sepconv1_act (Activatio	(None, 16, 32, 728)	0	add_10[0][0]
block13_sepconv1 (SeparableConv	(None, 16, 32, 728)	536536	block13_sepconv1_act[0][0]
block13_sepconv1_bn (BatchNorma	(None, 16, 32, 728)	2912	block13_sepconv1[0][0]
block13_sepconv2_act (Activatio	(None, 16, 32, 728)	0	block13_sepconv1_bn[0][0]

block13_sepconv2 (SeparableConv	(None, 16, 32, 1024)	752024	block13_sepconv2_act[0][0]
block13_sepconv2_bn (BatchNorma	(None, 16, 32, 1024)	4096	block13_sepconv2[0][0]
conv2d_3 (Conv2D)	(None, 8, 16, 1024)	745472	add_10[0][0]
block13_pool (MaxPooling2D)	(None, 8, 16, 1024)	0	block13_sepconv2_bn[0][0]
batch_normalization_3 (BatchNor	(None, 8, 16, 1024)	4096	conv2d_3[0][0]
add_11 (Add)	(None, 8, 16, 1024)	0	block13_pool[0][0] batch_normalization_3[0][0]
block14_sepconv1 (SeparableConv	(None, 8, 16, 1536)	1582080	add_11[0][0]
block14_sepconv1_bn (BatchNorma	(None, 8, 16, 1536)	6144	block14_sepconv1[0][0]
block14_sepconv1_act (Activatio	(None, 8, 16, 1536)	0	block14_sepconv1_bn[0][0]
block14_sepconv2 (SeparableConv	(None, 8, 16, 2048)	3159552	block14_sepconv1_act[0][0]
block14_sepconv2_bn (BatchNorma	(None, 8, 16, 2048)	8192	block14_sepconv2[0][0]
block14_sepconv2_act (Activatio	(None, 8, 16, 2048)	0	block14_sepconv2_bn[0][0]
global_average_pooling2d (Globa	(None, 2048)	0	block14_sepconv2_act[0][0]
dense (Dense)	(None, 1024)	2098176	global_average_pooling2d[0][0]
batch_normalization_4 (BatchNor	(None, 1024)	4096	dense[0][0]
dropout (Dropout)	(None, 1024)	0	batch_normalization_4[0][0]
dense_1 (Dense)	(None, 512)	524800	dropout[0][0]
batch_normalization_5 (BatchNor	(None, 512)	2048	dense_1[0][0]
dropout_1 (Dropout)	(None, 512)	0	batch_normalization_5[0][0]
dense_2 (Dense)	(None, 64)	32832	dropout_1[0][0]
dense_3 (Dense)	(None, 4)	260	dense_2[0][0]

=====

Total params: 23,523,692

Trainable params: 2,659,140

Non-trainable params: 20,864,552

```
In [ ]: log_dir=os.path.join("logs",datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
        tensorboard=tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1,write_graph=True,write_
        _grads=True)

        checkpoint_filepath='/content/drive//My Drive/Steel_Detection /multi_label.h5'
        model_checkpoint_callback=tf.keras.callbacks.ModelCheckpoint(filepath=checkpoint_filepath,monitor=
        'val_f1_score',mode='max',save_best_only=True)
        #https://keras.io/api/metrics/
        #https://keras.io/api/losses/probabilistic_losses/#categorical_crossentropy-function
        model.compile(optimizer='Adam',loss='binary_crossentropy',metrics=["acc",f1_score])
        callback=[model_checkpoint_callback,tensorboard]
        #https://datascience.stackexchange.com/questions/34444/what-is-the-difference-between-fit-and-fit-
        generator-in-keras
        history=model.fit_generator(train_image_generator,validation_data=val_image_generator,epochs=20,ve
        rbose=1,callbacks=callback)
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.

Epoch 1/20

150/150 [=====] - 2474s 16s/step - loss: 0.3016 - acc: 0.7647 - f1_score: 0.7640 - val_loss: 0.2212 - val_acc: 0.8008 - val_f1_score: 0.8246

Epoch 2/20

150/150 [=====] - 178s 1s/step - loss: 0.2332 - acc: 0.8087 - f1_score: 0.8120 - val_loss: 0.1843 - val_acc: 0.8450 - val_f1_score: 0.8628

Epoch 3/20

150/150 [=====] - 182s 1s/step - loss: 0.2066 - acc: 0.8310 - f1_score: 0.8408 - val_loss: 0.1967 - val_acc: 0.8358 - val_f1_score: 0.8525

Epoch 4/20

150/150 [=====] - 179s 1s/step - loss: 0.1998 - acc: 0.8362 - f1_score: 0.8460 - val_loss: 0.1849 - val_acc: 0.8267 - val_f1_score: 0.8536

Epoch 5/20

150/150 [=====] - 176s 1s/step - loss: 0.1869 - acc: 0.8441 - f1_score: 0.8530 - val_loss: 0.2130 - val_acc: 0.8183 - val_f1_score: 0.8384

Epoch 6/20

150/150 [=====] - 180s 1s/step - loss: 0.1848 - acc: 0.8456 - f1_score: 0.8537 - val_loss: 0.1733 - val_acc: 0.8450 - val_f1_score: 0.8705

Epoch 7/20

150/150 [=====] - 181s 1s/step - loss: 0.1779 - acc: 0.8489 - f1_score: 0.8590 - val_loss: 0.1759 - val_acc: 0.8375 - val_f1_score: 0.8610

Epoch 8/20

150/150 [=====] - 176s 1s/step - loss: 0.1756 - acc: 0.8543 - f1_score: 0.8640 - val_loss: 0.1740 - val_acc: 0.8375 - val_f1_score: 0.8614

Epoch 9/20

150/150 [=====] - 175s 1s/step - loss: 0.1681 - acc: 0.8521 - f1_score: 0.8641 - val_loss: 0.1634 - val_acc: 0.8500 - val_f1_score: 0.8746

Epoch 10/20

150/150 [=====] - 176s 1s/step - loss: 0.1682 - acc: 0.8546 - f1_score: 0.8663 - val_loss: 0.1592 - val_acc: 0.8583 - val_f1_score: 0.8800

Epoch 11/20

150/150 [=====] - 175s 1s/step - loss: 0.1604 - acc: 0.8558 - f1_score: 0.8698 - val_loss: 0.1621 - val_acc: 0.8542 - val_f1_score: 0.8758

Epoch 12/20

150/150 [=====] - 175s 1s/step - loss: 0.1645 - acc: 0.8554 - f1_score: 0.8683 - val_loss: 0.1766 - val_acc: 0.8275 - val_f1_score: 0.8558

Epoch 13/20

150/150 [=====] - 175s 1s/step - loss: 0.1657 - acc: 0.8521 - f1_score: 0.8677 - val_loss: 0.1596 - val_acc: 0.8508 - val_f1_score: 0.8727

Epoch 14/20

150/150 [=====] - 176s 1s/step - loss: 0.1546 - acc: 0.8704 - f1_score: 0.8794 - val_loss: 0.1585 - val_acc: 0.8642 - val_f1_score: 0.8866


```
Epoch 15/20
150/150 [=====] - 174s 1s/step - loss: 0.1544 - acc: 0.8598 - f1_score: 0.8757 - val_loss: 0.1545 - val_acc: 0.85
33 - val_f1_score: 0.8820
Epoch 16/20
150/150 [=====] - 174s 1s/step - loss: 0.1536 - acc: 0.8668 - f1_score: 0.8761 - val_loss: 0.1618 - val_acc: 0.85
92 - val_f1_score: 0.8823
Epoch 17/20
150/150 [=====] - 175s 1s/step - loss: 0.1532 - acc: 0.8664 - f1_score: 0.8790 - val_loss: 0.1546 - val_acc: 0.86
00 - val_f1_score: 0.8816
Epoch 18/20
150/150 [=====] - 174s 1s/step - loss: 0.1516 - acc: 0.8656 - f1_score: 0.8764 - val_loss: 0.1546 - val_acc: 0.86
00 - val_f1_score: 0.8789
Epoch 19/20
150/150 [=====] - 174s 1s/step - loss: 0.1504 - acc: 0.8664 - f1_score: 0.8837 - val_loss: 0.1552 - val_acc: 0.86
33 - val_f1_score: 0.8834
Epoch 20/20
150/150 [=====] - 174s 1s/step - loss: 0.1448 - acc: 0.8735 - f1_score: 0.8849 - val_loss: 0.1497 - val_acc: 0.86
75 - val_f1_score: 0.8841
```

```
In [ ]: %tensorboard --logdir logs
```

```
In [ ]: train_image_generator=val_datagen.flow_from_dataframe(dataframe=x_train,
                                                             directory=train_folder_path,
                                                             x_col="image_id",
                                                             y_col=col,
                                                             batch_size=32,
                                                             shuffle=False,
                                                             class_mode="raw",
                                                             target_size=(256,512))

val_image_generator=val_datagen.flow_from_dataframe(dataframe=x_val,
                                                    directory=train_folder_path,
                                                    x_col="image_id",
                                                    y_col=col,
                                                    batch_size=32,
                                                    shuffle=False,
                                                    class_mode="raw",
                                                    target_size=(256,512))

test_image_generator=val_datagen.flow_from_dataframe(dataframe=x_test,
                                                    directory=train_folder_path,
                                                    x_col="image_id",
                                                    y_col=col,
                                                    batch_size=32,
                                                    shuffle=False,
                                                    class_mode="raw",
                                                    target_size=(256,512))
```

```
Found 4799 validated image filenames.
Found 1200 validated image filenames.
Found 667 validated image filenames.
```

```
In [ ]: print('Training Dataset:\n')
print(model.evaluate(train_image_generator,verbose=1))
print("="*100)
print('\nValidation Dataset:\n')
print(model.evaluate(val_image_generator,verbose=1))
print("="*100)
print('\nTest Dataset:\n')
print(model.evaluate(test_image_generator,verbose=1))
```

Training Dataset:

150/150 [=====] - 52s 348ms/step - loss: 0.1153 - acc: 0.8960 - f1_score: 0.9109
[0.11531050503253937, 0.8960199952125549, 0.9108885526657104]

Validation Dataset:

38/38 [=====] - 13s 342ms/step - loss: 0.1497 - acc: 0.8675 - f1_score: 0.8840
[0.14966358244419098, 0.8675000071525574, 0.8840486407279968]

Test Dataset:

21/21 [=====] - 241s 12s/step - loss: 0.1459 - acc: 0.8711 - f1_score: 0.8915
[0.14590173959732056, 0.8710644841194153, 0.891469419002533]

- For validation, test datasets loss increases as compared to train but still as one move from validation to test slight decrease in loss and similar with metrics where validation, test datasets accuracy, f1_score decreases as compared to train but still as one move from validation to test increase in accuracy, f1_score can be observed which shows model works better on unseen data.

```
In [ ]:
```